

Prioritization strategies vs random (Target=fingpt, Dataset=FinSen)

Y: cumulative ANY bias hits in fingpt

% of inputs processed (effort)

Cumulative # of bias-revealing examples (fingpt)

20 40 60 80 100

10000

random

↓JSD (ref=distilroberta)

any-first (ref=distilroberta)

atomic-first (ref=distilroberta)

type-first (race→body→gender)

↓ length (longer first)

category-first (PropN→adj→noun→pronoun→verb)

8000

6000

4000

2000

0

0

20000

40000

60000

80000

Examples sampled

Milestones (share of final hits):				
random:	20=20%	40=40%	60=60%	80=80%
↓JSD (ref=distilroberta):	20=8%	40=34%	60=62%	80=83%
any-first (ref=distilroberta):	20=20%	40=40%	60=60%	80=80%
atomic-first (ref=distilroberta):	20=20%	40=40%	60=60%	80=80%
type-first (race→body→gender):	20=20%	40=41%	60=61%	80=81%
↓ length (longer first):	20=40%	40=45%	60=50%	80=82%
category-first (PropN→adj→noun→pronoun→verb):	20=19%	40=38%	60=58%	80=79%

