


Prioritization strategies vs random (Target=finma, Dataset=FinSen)

Y: cumulative ANY bias hits in finma

% of inputs processed (effort)

Cumulative # of bias-revealing examples (finma)



Milestones (share of final hits):

random: 20=20%, 40=40%, 60=60%, 80=80%

↓JSD (ref=distilroberta): 20=73%, 40=90%, 60=95%, 80=98%

any-first (ref=distilroberta): 20=23%, 40=42%, 60=62%, 80=81%

atomic-first (ref=distilroberta): 20=20%, 40=40%, 60=60%, 80=80%

type-first (body→gender→race): 20=42%, 40=57%, 60=71%, 80=86%

↓length (longer first): 20=32%, 40=80%, 60=86%, 80=88%

category-first (adj→noun→propN→verb→pron): 20=22%, 40=44%, 60=63%, 80=81%